

```
!pip install kaggle
```

```
Requirement already satisfied: kaggle in /usr/local/lib/python3.10/dist-packages (1.5.12)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.10/dist-packages (1.16.0)
Requirement already satisfied: certifi in /usr/local/lib/python3.10/dist-packages (2022.12.7)
Requirement already satisfied: python-dateutil in /usr/local/lib/python3.10/dist-packages (2.8.2)
Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (2.28.1)
Requirement already satisfied: tqdm in /usr/local/lib/python3.10/dist-packages (4.64.1)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.10/dist-packages (6.0.0)
Requirement already satisfied: urllib3 in /usr/local/lib/python3.10/dist-packages (1.26.15)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (5.0.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (0.11)
Requirement already satisfied: text-unidecode>=1.3 in /usr/local/lib/python3.10/dist-packages (1.3)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (3.1.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (3.4)
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount(
```

```
from logging import warning
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
df = pd.read_csv('/content/drive/MyDrive/Colab Notebooks/kaggle_dataset/customer_churn/Chu
```

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Balance
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	83803.61
2	3	15619304	Onio	502	France	Female	42	8	159660.88
3	4	15701354	Boni	699	France	Female	39	1	
4	5	15737888	Mitchell	850	Spain	Female	43	2	125560.85

```
df.columns
```

```
Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
      'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
```

```
'IsActiveMember', 'EstimatedSalary', 'Exited'],
dtype='object')
```

```
print('df:', df.shape)
```

```
df: (10000, 14)
```

```
df.duplicated().sum()
```

```
0
```

```
df.isnull().sum()
```

```
RowNumber      0
CustomerId      0
Surname         0
CreditScore     0
Geography       0
Gender          0
Age             0
Tenure          0
Balance         0
NumOfProducts  0
HasCrCard       0
IsActiveMember  0
EstimatedSalary 0
Exited          0
dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender                10000 non-null  object
6   Age                   10000 non-null  int64
7   Tenure                10000 non-null  int64
8   Balance               10000 non-null  float64
9   NumOfProducts         10000 non-null  int64
10  HasCrCard             10000 non-null  int64
11  IsActiveMember        10000 non-null  int64
12  EstimatedSalary       10000 non-null  float64
13  Exited                10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Balan
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.0000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.8892
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.4052
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.0000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.0000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.5400
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.2400
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.0900

```
plt.figure(figsize=(20,10))
sns.heatmap(df.corr(), annot = True)
```

<Axes: >

RowNumber	1	0.0042	0.0058	0.00078	-0.0065	-0.0091	0.0072	0.0006	0.012	-0.006
CustomerId	0.0042	1	0.0053	0.0095	-0.015	-0.012	0.017	-0.014	0.0017	0.015

```
col = []
for i in df.columns:
    num = len(df[i].unique())
    print(i,':', str(num) + str(' Distinct values'))
    #append < 10 distinct values columns in list
    if num < 10:
        col.append(i)
```

```
RowNumber : 10000 Distinct values
CustomerId : 10000 Distinct values
Surname : 2932 Distinct values
CreditScore : 460 Distinct values
Geography : 3 Distinct values
Gender : 2 Distinct values
Age : 70 Distinct values
Tenure : 11 Distinct values
Balance : 6382 Distinct values
NumOfProducts : 4 Distinct values
HasCrCard : 2 Distinct values
IsActiveMember : 2 Distinct values
EstimatedSalary : 9999 Distinct values
Exited : 2 Distinct values
```

```
for i in col:
    print(df[i].value_counts(), '\n')
```

```
France      5014
Germany     2509
Spain       2477
Name: Geography, dtype: int64
```

```
Male        5457
Female      4543
Name: Gender, dtype: int64
```

```
1      5084
2      4590
3       266
4        60
Name: NumOfProducts, dtype: int64
```

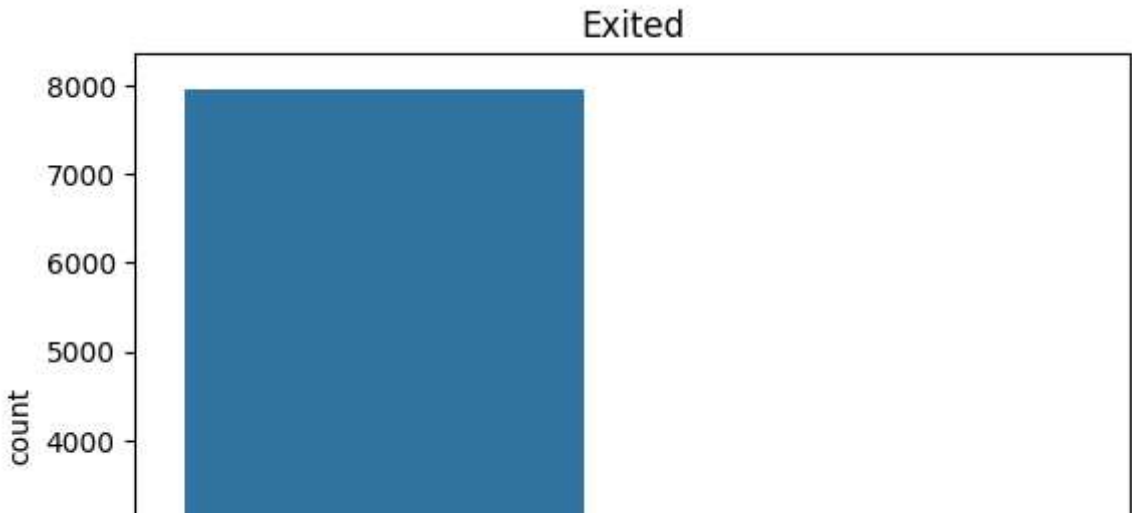
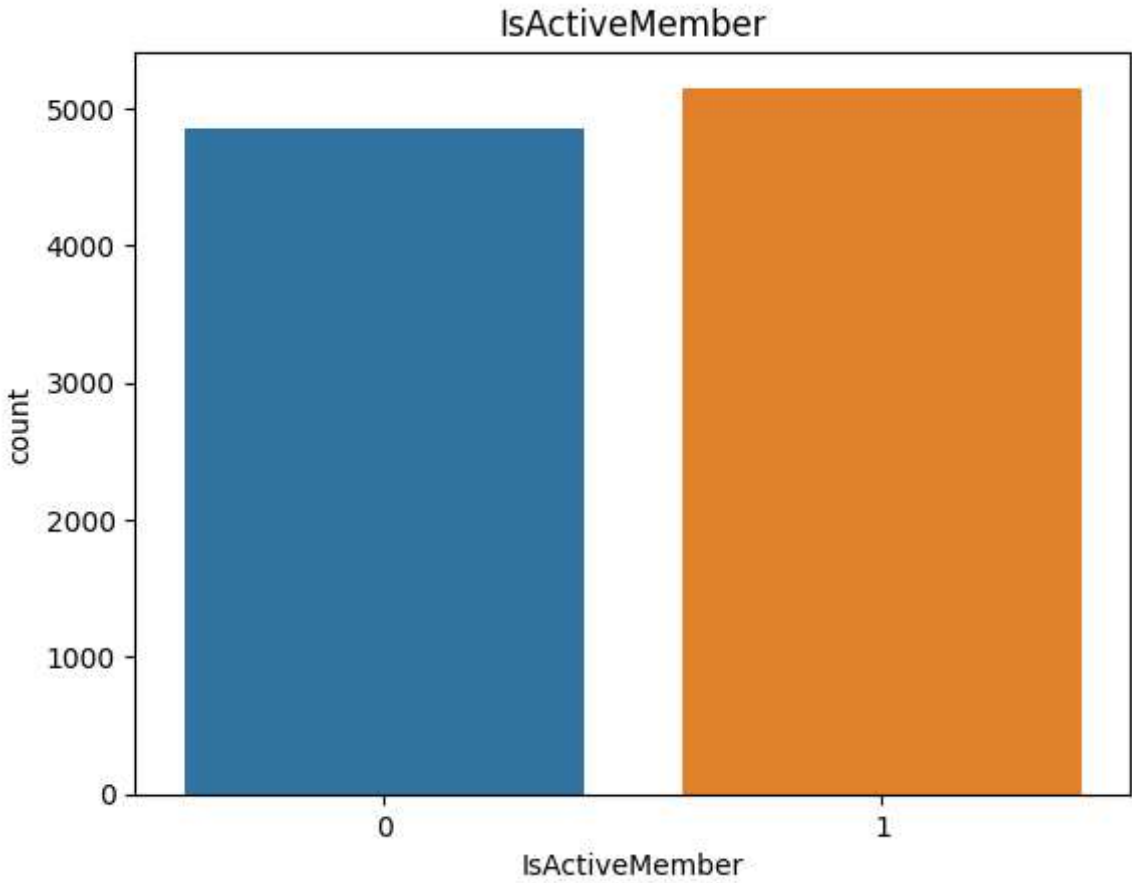

```
1      7055
0      2945
Name: HasCrCard, dtype: int64
```

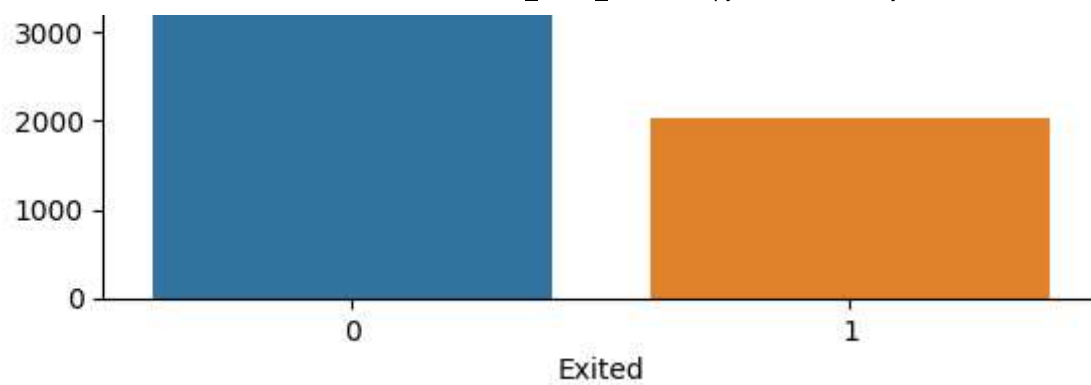
```
1      5151
0      4849
Name: IsActiveMember, dtype: int64
```

```
0      7963
1      2037
```


Name: Exited, dtype: int64

```
for i in col:  
    sns.countplot(x=i, data=df)  
    plt.title(i)  
    plt.show()
```

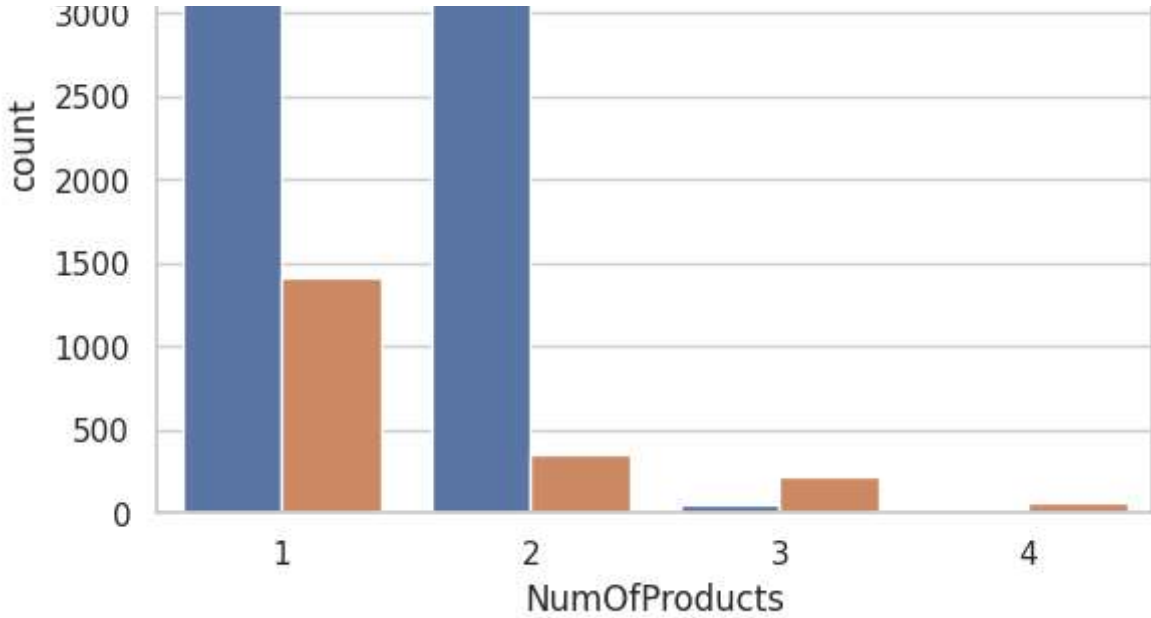




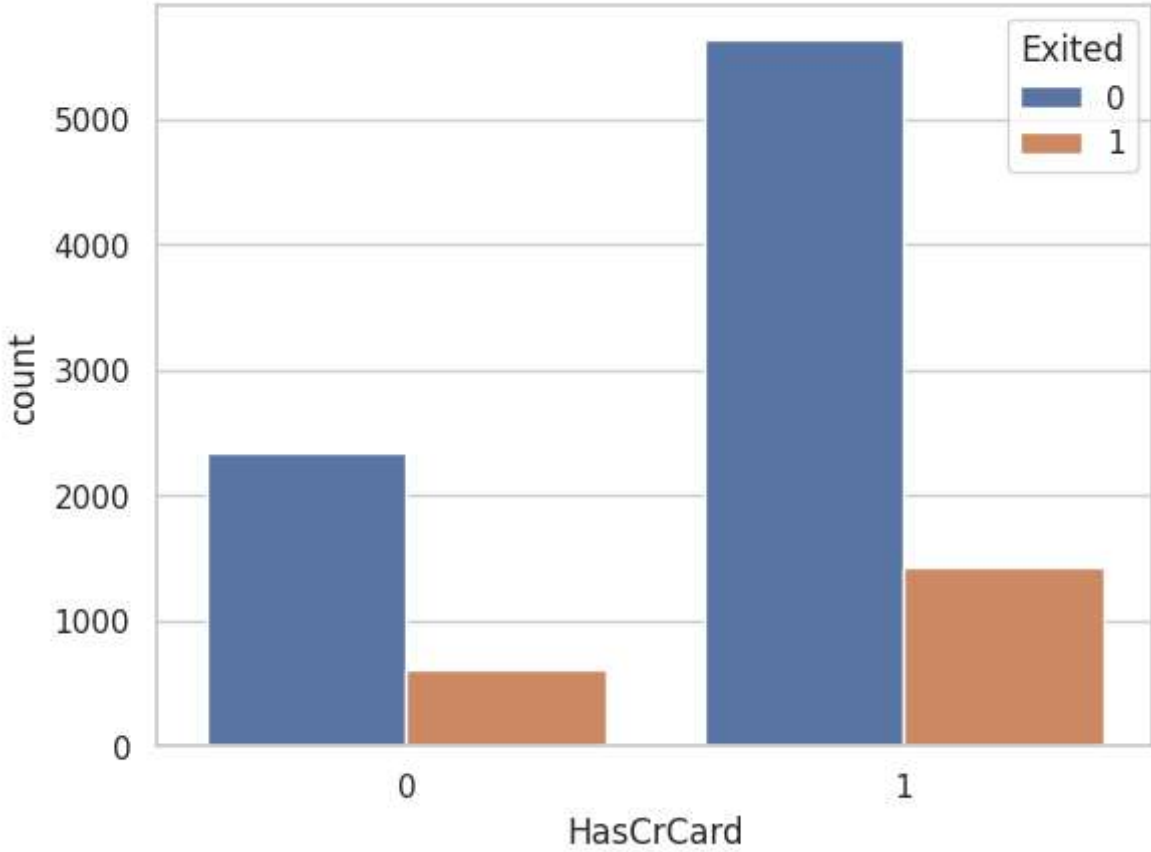
```
plt.hist(df['Age'], edgecolor='black')  
plt.title('Age Distribution')  
plt.show()
```




```
sns.set(style="whitegrid")
j = col.pop()
for i in col:
    sns.countplot(x=i, hue=j, data=df)
    plt.title( i + ' vs ' + j)
    plt.show()
```





HasCrCard vs Exited



IsActiveMember vs Exited





BLACKBOX AI

```
df.head()
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Bal
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	8380
2	3	15619304	Onio	502	France	Female	42	8	15960
3	4	15701354	Boni	699	France	Female	39	1	
4	5	15737888	Mitchell	850	Spain	Female	43	2	12550

```
df.drop(columns=['RowNumber', 'CustomerId', 'Surname'], inplace=True )
```

```
df.head()
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	
0	619	France	Female	42	2	0.00	1	1	
1	608	Spain	Female	41	1	83807.86	1	0	
2	502	France	Female	42	8	159660.80	3	1	
3	699	France	Female	39	1	0.00	2	0	
4	850	Spain	Female	43	2	125510.82	1	1	

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['Gender'] = le.fit_transform(df['Gender'])
```

```
df = pd.get_dummies(data = df, columns=['Geography'], drop_first=True)
```

```
df.head()
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActive
0	619	0	42	2	0.00	1	1	
1	608	0	41	1	83807.86	1	0	
2	502	0	42	8	159660.80	3	1	
3	699	0	39	1	0.00	2	0	
4	850	0	43	2	125510.82	1	1	

```
a = df.iloc[:, -3:-2]
```

```
a.head()
```

	Exited
0	1
1	0
2	1
3	0
4	0

```
df.drop(columns=['Exited'], inplace = True)
```

```
df.head()
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMeml
0	619	0	42	2	0.00	1	1	
1	608	0	41	1	83807.86	1	0	
2	502	0	42	8	159660.80	3	1	
3	699	0	39	1	0.00	2	0	
4	850	0	43	2	125510.82	1	1	

```
df = pd.concat([df, a], axis=1)
```

```
df.head()
```

	CreditScore	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMeml
0	619	0	42	2	0.00	1	1	
1	608	0	41	1	83807.86	1	0	
2	502	0	42	8	159660.80	3	1	
3	699	0	39	1	0.00	2	0	
4	850	0	43	2	125510.82	1	1	

```
plt.figure(figsize=(20,10))
sns.heatmap(df.corr(), annot = True)
```

<Axes: >

CreditScore	1	-0.0029	-0.004	0.00084	0.0063	0.012	-0.0055	0.026	-0.0014	0.0055	0.0048
Gender	-0.0029	1	-0.028	0.015	0.012	-0.022	0.0058	0.023	-0.0081	-0.025	0.017
Age	-0.004	-0.028	1	-0.01	0.028	-0.031	-0.012	0.085	-0.0072	0.047	-0.0017
Tenure	0.00084	0.015	-0.01	1	-0.012	0.013	0.023	-0.028	0.0078	-0.00057	0.0039
Balance	0.0063	0.012	0.028	-0.012	1	-0.3	-0.015	-0.01	0.013	0.4	-0.13
NumOfProducts	0.012	-0.022	-0.031	0.013	-0.3	1	0.0032	0.0096	0.014	-0.01	0.009
HasCrCard	-0.0055	0.0058	-0.012	0.023	-0.015	0.0032	1	-0.012	-0.0099	0.011	-0.013
IsActiveMember	0.026	0.023	0.085	-0.028	-0.01	0.0096	-0.012	1	-0.011	-0.02	0.017
EstimatedSalary	-0.0014	-0.0081	-0.0072	0.0078	0.013	0.014	-0.0099	-0.011	1	0.01	-0.0065
Geography_Germany	0.0055	-0.025	0.047	-0.00057	0.4	-0.01	0.011	-0.02	0.01	1	-0.33
Geography_Spain	0.0048	0.017	-0.0017	0.0039	-0.13	0.009	-0.013	0.017	-0.0065	-0.33	1
Exited	-0.027	-0.11	0.29	-0.014	0.12	-0.048	-0.0071	-0.16	0.012	0.17	-0.053
	itScore	Gender	Age	Tenure	Balance	Products	CrCard	Member	Salary	Germany	Spain

```
x = df.iloc[:, :-1].values
y = df.iloc[:, -1].values
```

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.2, random_state=
```

```
print(x_train.shape, x_test.shape)
print(y_train.shape, y_test.shape)
```

```
(8000, 11) (2000, 11)
(8000,) (2000,)
```

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
```

```
from sklearn.linear_model import LogisticRegression
log = LogisticRegression(random_state = 42)
log.fit(x_train, y_train)
```

```
LogisticRegression
LogisticRegression(random_state=42)
```

```
from sklearn.ensemble import GradientBoostingClassifier, RandomForestClassifier
R = RandomForestClassifier(n_estimators=5 ,random_state = 42)
```

```
R.fit(x_train, y_train)
```

```
▼ RandomForestClassifier  
RandomForestClassifier(n_estimators=5, random_state=42)
```


```
clf = GradientBoostingClassifier(n_estimators=10, learning_rate=1.0, random_state=42)  
clf.fit(x_train, y_train)
```

```
▼ GradientBoostingClassifier  
GradientBoostingClassifier(learning_rate=1.0, n_estimators=10, random_state=42)
```

```
classifier = [log, R, clf]  
model = ['Logistic Regression', 'Random Forest Classifier', 'Gradient Boosting Classifier']
```

```
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score, Confu  
for i in range(len(classifier)):  
    y_pred = classifier[i].predict(x_test)  
    cm = confusion_matrix(y_test, y_pred)  
    accuracy = accuracy_score(y_test, y_pred)*100  
    print('\nfor ' + str(model[i]) + ':\n')  
    disp = ConfusionMatrixDisplay(confusion_matrix=cm)  
    plt.rcParams['axes.grid'] = False  
    disp.plot()  
    print(accuracy)  
    print(classification_report(y_test, y_pred))  
    plt.show()
```

	0	0.88	0.92	0.90	1607
	1	0.59	0.48	0.53	393
accuracy				0.83	2000
macro avg		0.73	0.70	0.71	2000
weighted avg		0.82	0.83	0.83	2000



for Gradient Boosting Classifier:

85.65

	precision	recall	f1-score	support
0	0.88	0.95	0.91	1607
1	0.70	0.47	0.56	393
accuracy			0.86	2000
macro avg	0.79	0.71	0.74	2000
weighted avg	0.84	0.86	0.84	2000

